

$$\hat{\beta}_{Ridge} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$$

with  $\mathbf{I}$  being a  $p \times p$  identity matrix.

The ordinary least squares result is

$$\hat{\beta}_{OLS} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

To show that the optimal parameters for Ridge regression are given by:

$$\hat{\beta}_{Ridge} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$$

we need to derive this from the Ridge regression cost function.

The cost function for Ridge regression is given by:

$$C(\beta) = (\mathbf{y} - \mathbf{X}\beta)^T (\mathbf{y} - \mathbf{X}\beta) + \lambda \beta^T \beta$$

Expand the Cost Function:

$$C(\beta) = (\mathbf{y}^T - \beta^T \mathbf{X}^T)(\mathbf{y} - \mathbf{X}\beta) + \lambda \beta^T \beta$$

$$C(\beta) = \mathbf{y}^T \mathbf{y} - \mathbf{y}^T \mathbf{X}\beta - \beta^T \mathbf{X}^T \mathbf{y} + \beta^T \mathbf{X}^T \mathbf{X}\beta + \lambda \beta^T \beta$$

then combine the terms:

$$C(\beta) = \mathbf{y}^T \mathbf{y} - 2\mathbf{y}^T \mathbf{X}\beta + \beta^T (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})\beta$$

Take the Gradient with Respect to  $\beta$  and set equal to zero for optimization:

$$\frac{\partial C(\beta)}{\partial \beta} = -2\mathbf{X}^T \mathbf{y} + 2(\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})\beta = 0$$

$$-2\mathbf{X}^T \mathbf{y} + 2(\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})\beta = 0$$

$$(\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})\beta = \mathbf{X}^T \mathbf{y}$$

Solve for  $(\beta)$ :

$$\beta = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$$

Thus, we have shown that the optimal parameters for Ridge regression are:

$\mathbf{U}$  and  $\mathbf{V}$  are orthogonal matrices of dimensions  $n \times n$  and  $p \times p$ , and  $\mathbf{\Sigma}$  is an  $n \times p$  matrix which contains the singular values only.

First we show that you can write the OLS solutions in terms of the eigenvectors (the columns) of the orthogonal matrix  $\mathbf{U}$  as

$$\tilde{\mathbf{y}}_{OLS} = \mathbf{X}\beta = \sum_{j=0}^{p-1} \mathbf{u}_j \mathbf{u}_j^T \mathbf{y}$$

The OLS solution for  $\beta$  is given by:

$$\hat{\beta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

First we can find  $\mathbf{X}^T \mathbf{X}$ :

$$\mathbf{X}^T \mathbf{X} = (\mathbf{U} \mathbf{\Sigma} \mathbf{V}^T)^T (\mathbf{U} \mathbf{\Sigma} \mathbf{V}^T) = \mathbf{V} \mathbf{\Sigma}^T \mathbf{U}^T \mathbf{U} \mathbf{\Sigma} \mathbf{V}^T = \mathbf{V} \mathbf{\Sigma}^T \mathbf{\Sigma} \mathbf{V}^T$$

Since  $\mathbf{U}$  is orthogonal,  $\mathbf{U}^T \mathbf{U} = \mathbf{I}$ .

Inverse of  $\mathbf{X}^T \mathbf{X}$ :

$$(\mathbf{X}^T \mathbf{X})^{-1} = (\mathbf{V} \mathbf{\Sigma}^T \mathbf{\Sigma} \mathbf{V}^T)^{-1} = \mathbf{V}^{-T} (\mathbf{\Sigma}^T \mathbf{\Sigma})^{-1} \mathbf{V}^{-1} = \mathbf{V} (\mathbf{\Sigma}^T \mathbf{\Sigma})^{-1} \mathbf{V}^T$$

Since  $\mathbf{V}$  is orthogonal,  $\mathbf{V}^{-1} = \mathbf{V}^T$ .

Substitute Back into OLS Solution:

$$\hat{\beta} = \mathbf{V} (\mathbf{\Sigma}^T \mathbf{\Sigma})^{-1} \mathbf{V}^T \mathbf{X}^T \mathbf{y}$$

$$\hat{\beta} = \mathbf{V} (\mathbf{\Sigma}^T \mathbf{\Sigma})^{-1} \mathbf{V}^T \mathbf{V} \mathbf{\Sigma}^T \mathbf{U}^T \mathbf{y}$$

$$\hat{\beta} = \mathbf{V} (\mathbf{\Sigma}^T \mathbf{\Sigma})^{-1} \mathbf{\Sigma}^T \mathbf{U}^T \mathbf{y}$$

Compute  $\tilde{\mathbf{y}}_{OLS}$ :

$$\tilde{\mathbf{y}}_{OLS} = \mathbf{X} \hat{\beta} = \mathbf{U} \mathbf{\Sigma} \mathbf{V}^T \mathbf{V} (\mathbf{\Sigma}^T \mathbf{\Sigma})^{-1} \mathbf{\Sigma}^T \mathbf{U}^T \mathbf{y}$$

$$\tilde{\mathbf{y}}_{OLS} = \mathbf{U} \mathbf{\Sigma} (\mathbf{\Sigma}^T \mathbf{\Sigma})^{-1} \mathbf{\Sigma}^T \mathbf{U}^T \mathbf{y}$$

And since  $(\mathbf{\Sigma}^T \mathbf{\Sigma}) = \text{diag}(\sigma_0^2, \sigma_1^2, \dots, \sigma_i^2)$

and  $(\mathbf{\Sigma}^T \mathbf{\Sigma})^{-1} = \text{diag}(1/\sigma_0^2, 1/\sigma_1^2, \dots, 1/\sigma_i^2)$  then  $\mathbf{\Sigma} (\mathbf{\Sigma}^T \mathbf{\Sigma})^{-1} \mathbf{\Sigma}^T = \mathbf{I}$ , so:

$$\tilde{\mathbf{y}}_{OLS} = \sum_{j=0}^{p-1} \mathbf{u}_j \mathbf{u}_j^T \mathbf{y}$$

## Ridge

For Ridge regression, show that the corresponding equation is

$$\tilde{\mathbf{y}}_{Ridge} = \mathbf{X} \beta_{Ridge} = \mathbf{U} \Sigma \mathbf{V}^T (\mathbf{V} \Sigma^2 \mathbf{V}^T + \lambda \mathbf{I})^{-1} (\mathbf{U} \Sigma \mathbf{V}^T)^T \mathbf{y} = \sum_{j=1}^{p-1} \mathbf{u}_j \mathbf{u}_j^T \frac{\sigma_j^2}{\sigma_j^2 + \lambda} \mathbf{y}$$

First we start with:

$$\mathbf{U} \Sigma \mathbf{V}^T (\mathbf{V} \Sigma^2 \mathbf{V}^T + \lambda \mathbf{I})^{-1} (\mathbf{U} \Sigma \mathbf{V}^T)^T \mathbf{y}$$

and re-write to:

$$\mathbf{U} \Sigma \mathbf{V}^T V (\Sigma^2 + \lambda \mathbf{I})^{-1} V^T (\mathbf{U} \Sigma \mathbf{V}^T)^T \mathbf{y}$$

Then we can further simplify:

$$\mathbf{U} \Sigma (\Sigma^2 + \lambda \mathbf{I})^{-1} \Sigma^T \mathbf{U}^T \mathbf{y}$$

Now we can take a look at  $\Sigma (\Sigma^2 + \lambda \mathbf{I})^{-1} \Sigma^T$  element by element,

$$\Sigma (\Sigma^2 + \lambda \mathbf{I})^{-1} \Sigma^T = \left( \frac{\sigma_0^2}{\sigma_0^2 + \lambda}, \frac{\sigma_1^2}{\sigma_1^2 + \lambda}, \frac{\sigma_i^2}{\sigma_i^2 + \lambda} \right)$$

Final expression:

$$\tilde{\mathbf{y}}_{Ridge} = \mathbf{U} \Sigma \mathbf{V}^T V (\Sigma^2 + \lambda \mathbf{I})^{-1} V^T (\mathbf{U} \Sigma \mathbf{V}^T)^T \mathbf{y} = \sum_{j=1}^{p-1} \mathbf{u}_j \mathbf{u}_j^T \frac{\sigma_j^2}{\sigma_j^2 + \lambda} \mathbf{y}$$

## Exercise 2: Adding Ridge Regression

```
In [ ]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
from sklearn.model_selection import train_test_split
```

```

In [ ]: X = np.ones((n,6))
X[:,1] = x[:,0]
X[:,2] = x[:,0]**2
X[:,3] = x[:,0]**3
X[:,4] = x[:,0]**4
X[:,5] = x[:,0]**5

def scaling(X,y):
    xmean = np.mean(X, axis=0)
    xstd = np.std(X, axis=0)
    xstd[xstd == 0] = 1 #avoid division by zero
    ymean = np.mean(y)
    ystd = np.std(y)
    X_scaled = (X - xmean)/xstd
    y_scaled = (y - ymean)/ystd
    X_scaled[:,0] = 1
    return X_scaled, y_scaled, xmean, xstd, ymean, ystd

X_scaled, y_scaled, xmean, xstd, ymean, ystd = scaling(X,y)
print(X_scaled.shape)

```

(100, 6)

```

In [ ]: #test train split
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y_scaled, t
sort = np.argsort(X_train[:,1])
X_train = X_train[sort]
y_train = y_train[sort]

sort = np.argsort(X_test[:,1])
X_test = X_test[sort]
y_test = y_test[sort]

```

```

In [ ]: #calculate the beta values for OLS and Ridge
lam = [0.0001, 0.001, 0.01, 0.1, 1]
betaOLS = []
betaRidge = []

for i in lam:
    ols = np.linalg.inv(X_train.T @ X_train) @ X_train.T @ y_train
    ridge = np.linalg.inv(X_train.T @ X_train + i*np.eye(6)) @ X_train.T

    betaOLS.append(ols)
    betaRidge.append(ridge)

```

```

In [ ]: for i in range(len(lam)):
        ytildeOLS = X_train @ betaOLS[i]

```

```

print('MSE for Ridge with lambda = %f' % mse_ridge)

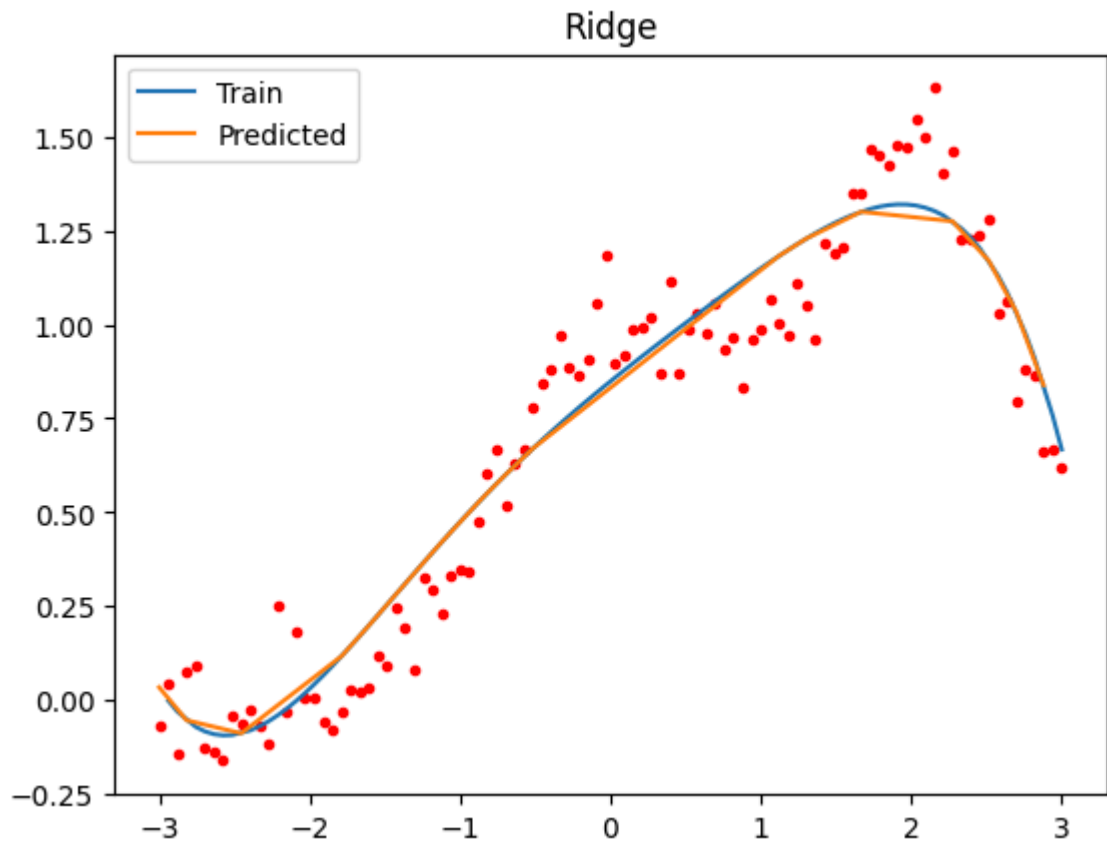
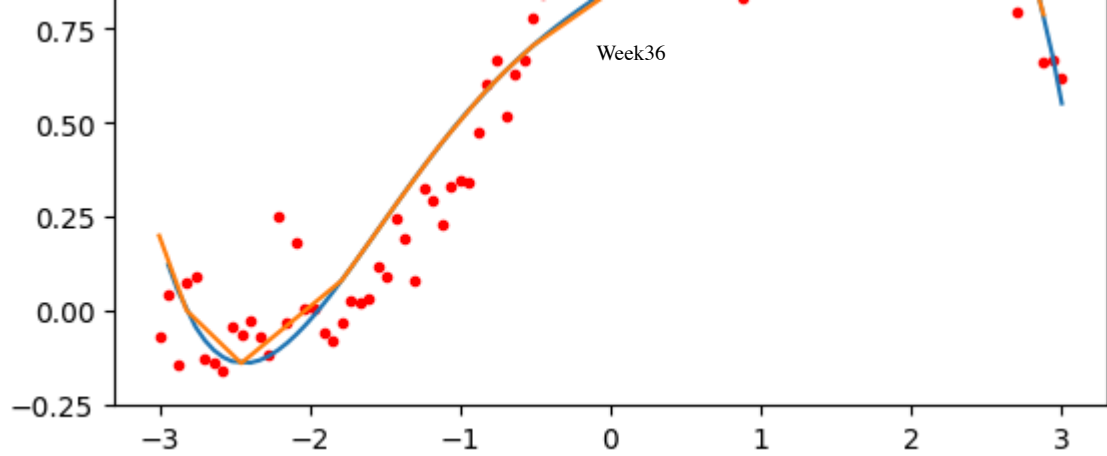
# #unscale
X_train = X_train*xstd + xmean
X_test = X_test*xstd + xmean
ytildeOLS = ytildeOLS*ystd + ymean
ytildeRidge = ytildeRidge*ystd + ymean
ypredictOLS = ypredictOLS*ystd + ymean
ypredictRidge = ypredictRidge*ystd + ymean

#plot the data
plt.plot(x, y, 'ro', markersize = 3)
plt.plot(X_train[:,1], ytildeOLS, label='Train')
plt.plot(X_test[:,1], ypredictOLS, label='Predicted')
plt.title('OLS')
plt.legend()
plt.show()

#plot the data
plt.plot(x, y, 'ro', markersize = 3)
plt.plot(X_train[:,1], ytildeRidge, label='Train')
plt.plot(X_test[:,1], ypredictRidge, label='Predicted')
plt.title('Ridge')
plt.legend()
plt.show()

```

MSE for OLS with	lambda	= None	is	test: 0.0689	train: 0.0640
MSE for Ridge with	lambda	= 0.0001	is	test: 0.0689	train: 0.0640
-----					
MSE for OLS with	lambda	= None	is	test: 0.0689	train: 0.0640
MSE for Ridge with	lambda	= 0.0010	is	test: 0.0688	train: 0.0640
-----					
MSE for OLS with	lambda	= None	is	test: 0.0689	train: 0.0640
MSE for Ridge with	lambda	= 0.0100	is	test: 0.0684	train: 0.0640
-----					
MSE for OLS with	lambda	= None	is	test: 0.0689	train: 0.0640
MSE for Ridge with	lambda	= 0.1000	is	test: 0.0655	train: 0.0644
-----					
MSE for OLS with	lambda	= None	is	test: 0.0689	train: 0.0640
MSE for Ridge with	lambda	= 1.0000	is	test: 0.0673	train: 0.0702
-----					



```
In [ ]: #now we do the same with 10 and 15 degrees
deg = [10, 15]
for d in deg:
    X = PolynomialFeatures(degree=d).fit_transform(x)
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0
```

```

sort = np.argsort(X_test[:,1])
X_test = X_test[sort]
y_test = y_test[sort]

betaOLS = []
betaRidge = []
lam = [0.0001, 0.001, 0.01, 0.1, 1]
for i in lam:
    ols = np.linalg.inv(X_train.T @ X_train) @ X_train.T @ y_train
    ridge = np.linalg.inv(X_train.T @ X_train + i*np.eye(d+1)) @ X_train.T @ y_train

    betaOLS.append(ols)
    betaRidge.append(ridge)

for i in range(len(lam)):
    ytildeOLS = X_train @ betaOLS[i]
    ytildeRidge = X_train @ betaRidge[i]

    ypredictOLS = X_test @ betaOLS[i]
    ypredictRidge = X_test @ betaRidge[i]

    MSEOLStest = mse(y_test, ypredictOLS)
    MSERidgetest = mse(y_test, ypredictRidge)
    MSEOLStrain = mse(y_train, ytildeOLS)
    MSERidgetrain = mse(y_train, ytildeRidge)

    print(f'MSE for OLS with lambda = None is |test: {MSEOLStest} |train: {MSEOLStrain}')
    print(f'MSE for Ridge with lambda = {lam[i]:.4f} is |test: {MSERidgetest} |train: {MSERidgetrain}')
    print('-----')

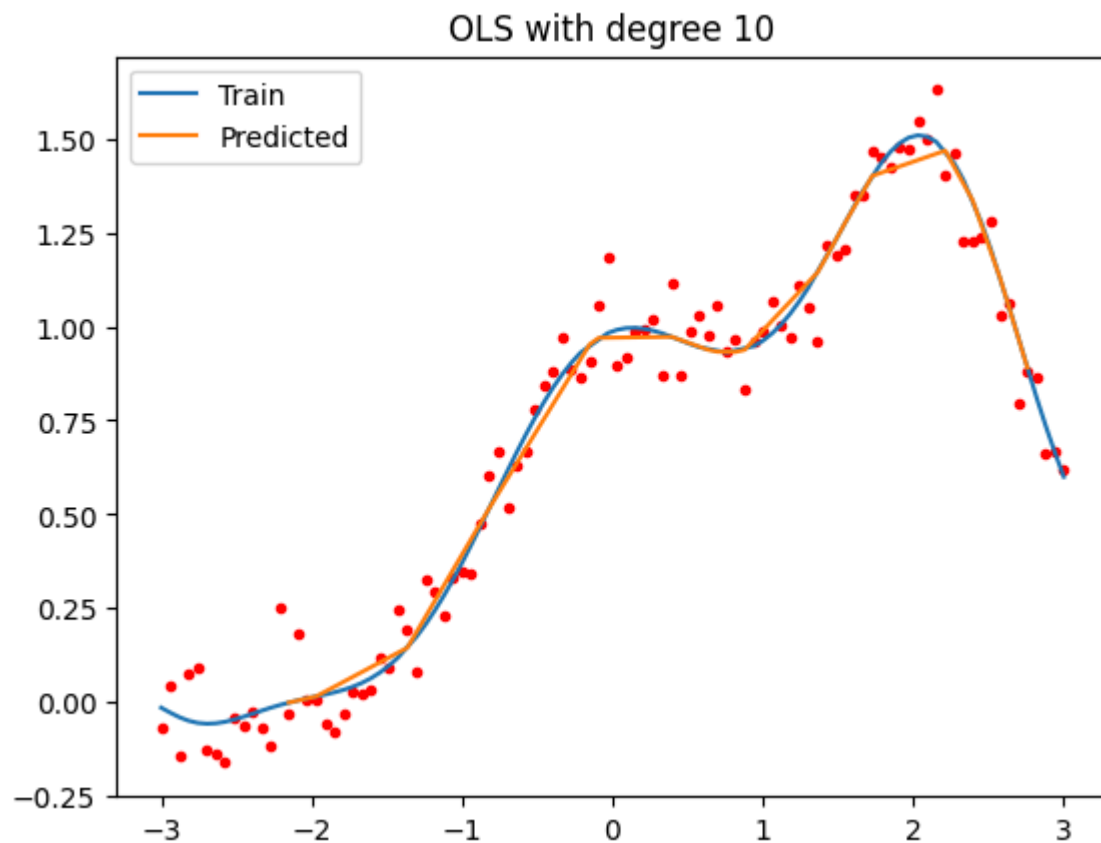
# #unscale for plotting
# for i in range(1, d+1):
#     X_train[:, i] = X_train[:, i] - means[i-1]
#     X_test[:, i] = X_test[:, i] - means[i-1]

#unscale using sklearn
X_train[:,1:] = scaler.inverse_transform(X_train[:,1:])
X_test[:,1:] = scaler.inverse_transform(X_test[:,1:])

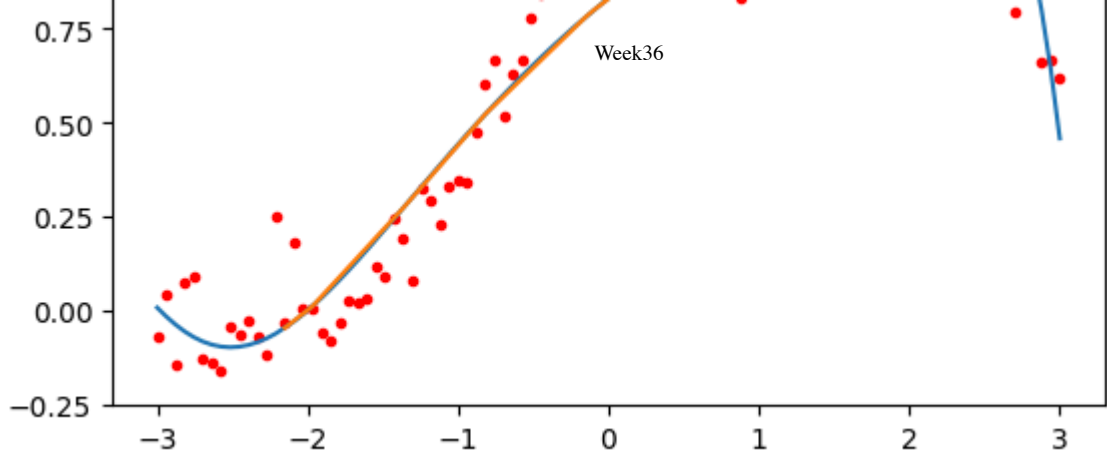
#plot the data
plt.plot(x, y, 'ro', markersize = 3)
plt.plot(X_train[:,1], ytildeOLS, label='Train')
plt.plot(X_test[:,1], ypredictOLS, label='Predicted')
plt.title(f'OLS with degree {d}')
plt.legend()
plt.show()

```


MSE for Ridge with lambda = 0.0010	is	test: 0.0092	train: 0.0066
-----			
MSE for OLS with lambda = None	is	test: 0.0086	train: 0.0060
MSE for Ridge with lambda = 0.0100	is	test: 0.0118	train: 0.0076
-----			
MSE for OLS with lambda = None	is	test: 0.0086	train: 0.0060
MSE for Ridge with lambda = 0.1000	is	test: 0.0210	train: 0.0119
-----			
MSE for OLS with lambda = None	is	test: 0.0086	train: 0.0060
MSE for Ridge with lambda = 1.0000	is	test: 0.0239	train: 0.0154
-----			







MSE for OLS with	lambda	= None	is	test: 0.0055	train: 0.0063
MSE for Ridge with	lambda	= 0.0001	is	test: 0.0047	train: 0.0066
-----					
MSE for OLS with	lambda	= None	is	test: 0.0055	train: 0.0063
MSE for Ridge with	lambda	= 0.0010	is	test: 0.0055	train: 0.0069
-----					
MSE for OLS with	lambda	= None	is	test: 0.0055	train: 0.0063
MSE for Ridge with	lambda	= 0.0100	is	test: 0.0088	train: 0.0084
-----					
MSE for OLS with	lambda	= None	is	test: 0.0055	train: 0.0063
MSE for Ridge with	lambda	= 0.1000	is	test: 0.0151	train: 0.0109
-----					
MSE for OLS with	lambda	= None	is	test: 0.0055	train: 0.0063
MSE for Ridge with	lambda	= 1.0000	is	test: 0.0238	train: 0.0151
-----					



Ridge with degree 15

